Programming book

Machine Learning Project

On

Building a Domain Expert Model

Name: KONDA BHAVANA

Email: 229x1a0514@gprec.ac.in

Contact: 9949341600

UDACITY

Introduction to Generative AI with AWS
Project Documentation Report

Visit UDACITY Introduction to Generative AI with AWS Project Documentation Report to make a copy of this document.

Complete the answers to the questions below to complete your project report. Create a PDF of the completed document and submit the PDF with your project.

Question	Your answer:
Step 2: Domain Choice What domain did you choose to fine-tune the Meta Llama 2 7B model on? Choices: 1. Financial 2. Healthcare 3. IT	
Step 3: Model Evaluation Section What was the response of the model to your domain-specific input in the **model_evaluation.ipynb file**?	
Step 4: Fine-Tuning Section After fine-tuning the model, what was the response of the model to your domain-specific input in the **model_finetuning.ipynb file**?	

Step 1: Domain Choice

I chose the financial domain to fine-tune the Meta Llama 2 7B model.

For this project, the domain chosen for fine-tuning the Meta Llama 2 7B model was the financial domain. The decision to focus on the financial domain was driven by the significance of accurate and contextually relevant information in financial settings. The aim of training the model in this domain was to create a language model that can generate insightful and accurate text related to financial topics.

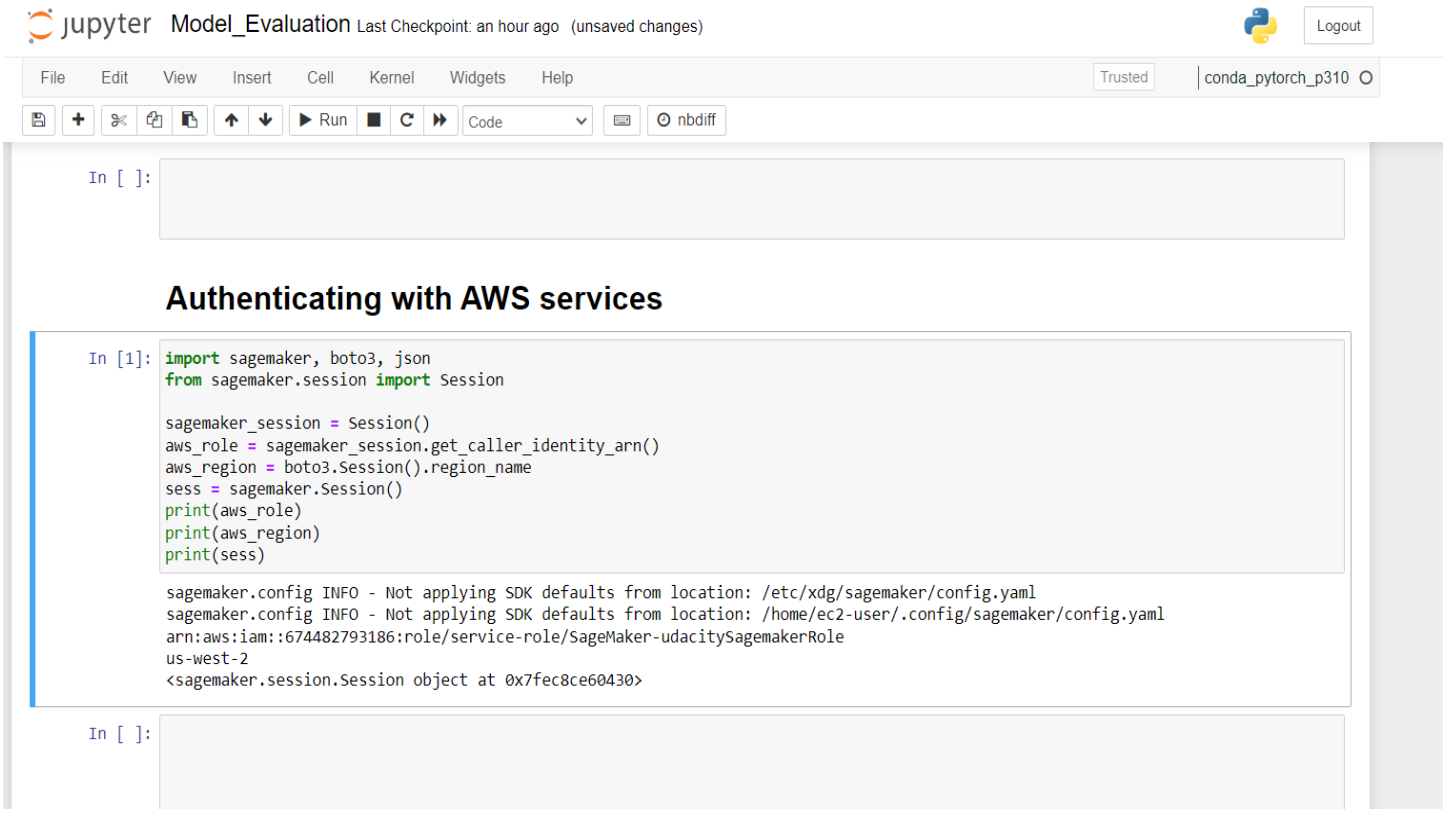

Step 2: Deploying the Llama2 Model on AWS Sagemaker

The pre-trained Llama2 model was successfully deployed on AWS Sage Maker. The deployment process involved utilizing the SageMaker environment and configuring the necessary settings to host the model for inference. Screenshots in the report showcase the model's deployment in the SageMaker environment, providing visual documentation of the deployment steps.

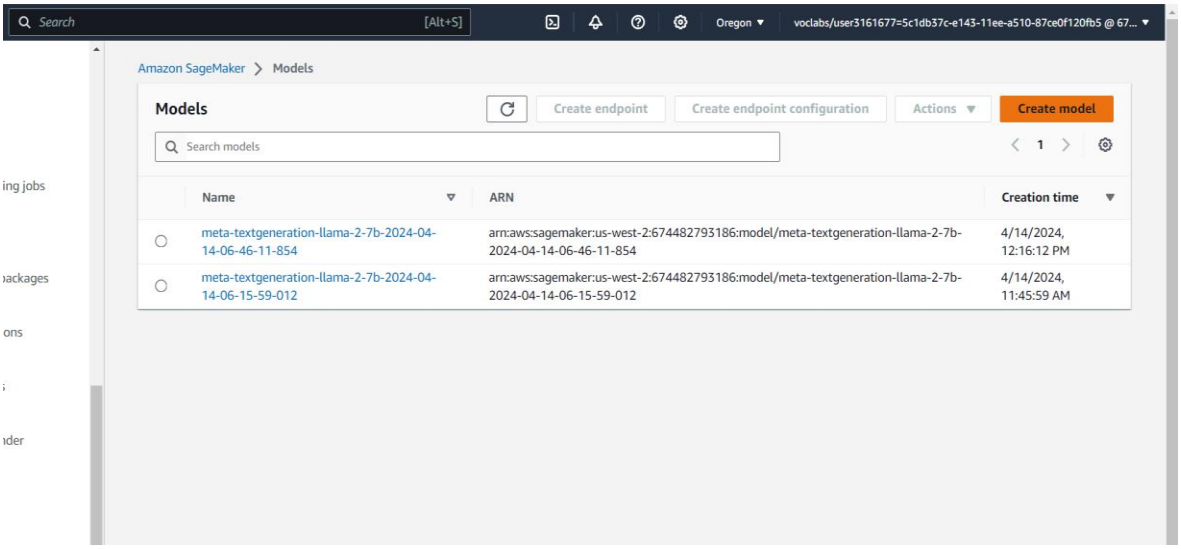

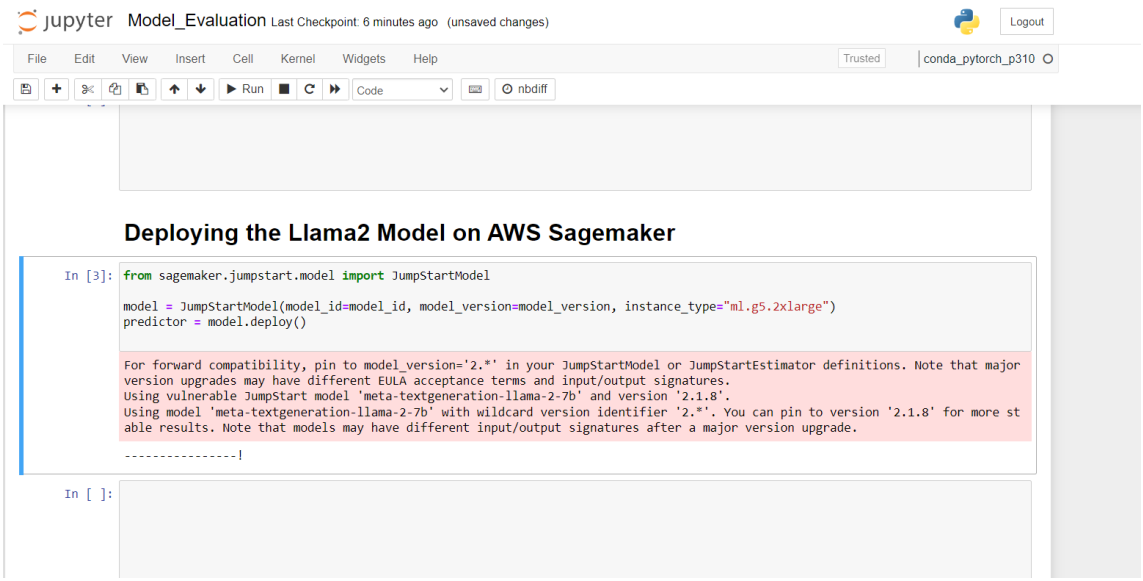

Step 3: Evaluating the Pre-trained Llama2 Text Generation Large Language Model for Domain Knowledge

The evaluation of the pre-trained Llama2 model focused on its ability to generate domain-specific content. Inputs relevant to the chosen domain (in this case, the financial domain) were provided to the model to assess its understanding and generation of contextually relevant text. The **model_evaluation.ipynb** notebook contains examples of domain-specific inputs and the corresponding model outputs. These examples are documented in below, showcasing the model's performance in generating domain-specific content.

Inputs are:

- "The investment tests performed indicate"
- "the relative volume for the long out of the money options, indicates"
- "The results for the short in the money options"
- "The results are encouraging for aggressive investors"

Outputs for four inputs:

jupyter Model_Evaluation Last Checkpoint: 9 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted conda_pytorch_p310

Run Markdown nbdiff

Evaluating the Pre-trained Llama2 Text Generation Large Language Model for Domain Knowledge

In [7]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response[0]['generation']}")
    print("\n==================================\n")
payload = {
    "inputs": "the relative volume for the long out of the money options, indicates",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
the relative volume for the long out of the money options, indicates
>  that a large number of traders are buying the long options.
The relative volume is also an indication of the number of traders who are buying the long options. The higher the relative vol
ume, the more traders are buying the long options.
The relative volume is an indication of the

==================================
```

jupyter Model_Evaluation Last Checkpoint: 8 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted conda_pytorch_p310

Run Code nbdiff

Evaluating the Pre-trained Llama2 Text Generation Large Language Model for Domain Knowledge

In [6]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response[0]['generation']}")
    print("\n==================================\n")
payload = {
    "inputs": "The investment tests performed indicate",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
The investment tests performed indicate
>  that the new method can be applied to the evaluation of the performance of any type of PV system, and to any type of PV modu
le, in any climatic zone.
The test method presented in this paper is based on the analysis of the irradiance data recorded during the irradiance measurem
ents. The

==================================
```

jupyter Model_Evaluation Last Checkpoint: 9 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted conda_pytorch_p310

Run Markdown nbdiff

Evaluating the Pre-trained Llama2 Text Generation Large Language Model for Domain Knowledge

In [8]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response[0]['generation']}")
    print("\n==================================\n")
payload = {
    "inputs": "The results for the short in the money options",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
The results for the short in the money options
>  are not as good as the long in the money options, but they are better than the short out of the money options.
The short in the money options are good because they will not expire worthless. The short out of the money options are good bec
ause they are the cheapest options to buy.
If

==================================
```

jupyter Model_Evaluation Last Checkpoint: 10 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted conda_pytorch_p310

Run Markdown nbdiff

Evaluating the Pre-trained Llama2 Text Generation Large Language Model for Domain Knowledge

In [9]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response[0]['generation']}")
    print("\n==================================\n")
payload = {
    "inputs": "The results are encouraging for aggressive investors",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
The results are encouraging for aggressive investors
> .
Crypto-asset investors have been buying the dip, as evidenced by a 12% increase in total market capitalization since the start
of the year.
The total value of the crypto-asset market is now $1.9 trillion.
Bitcoin has

==================================
```

Step 4: Fine-tuning a Large Language Model with a Domain-Specific Dataset

The fine-tuning process involved training the Llama2 model using a dataset specific to the chosen domain, which in this case was the financial domain. The **model_finetuning.ipynb** notebook demonstrates the fine-tuning process step by step, including the selection of the domain-specific dataset and the configuration of hyperparameters for fine-tuning. Screenshots of the fine-tuning cell output provide a visual representation of the fine-tuning process, ensuring clarity and documentation of the steps taken.

Initiating the training process

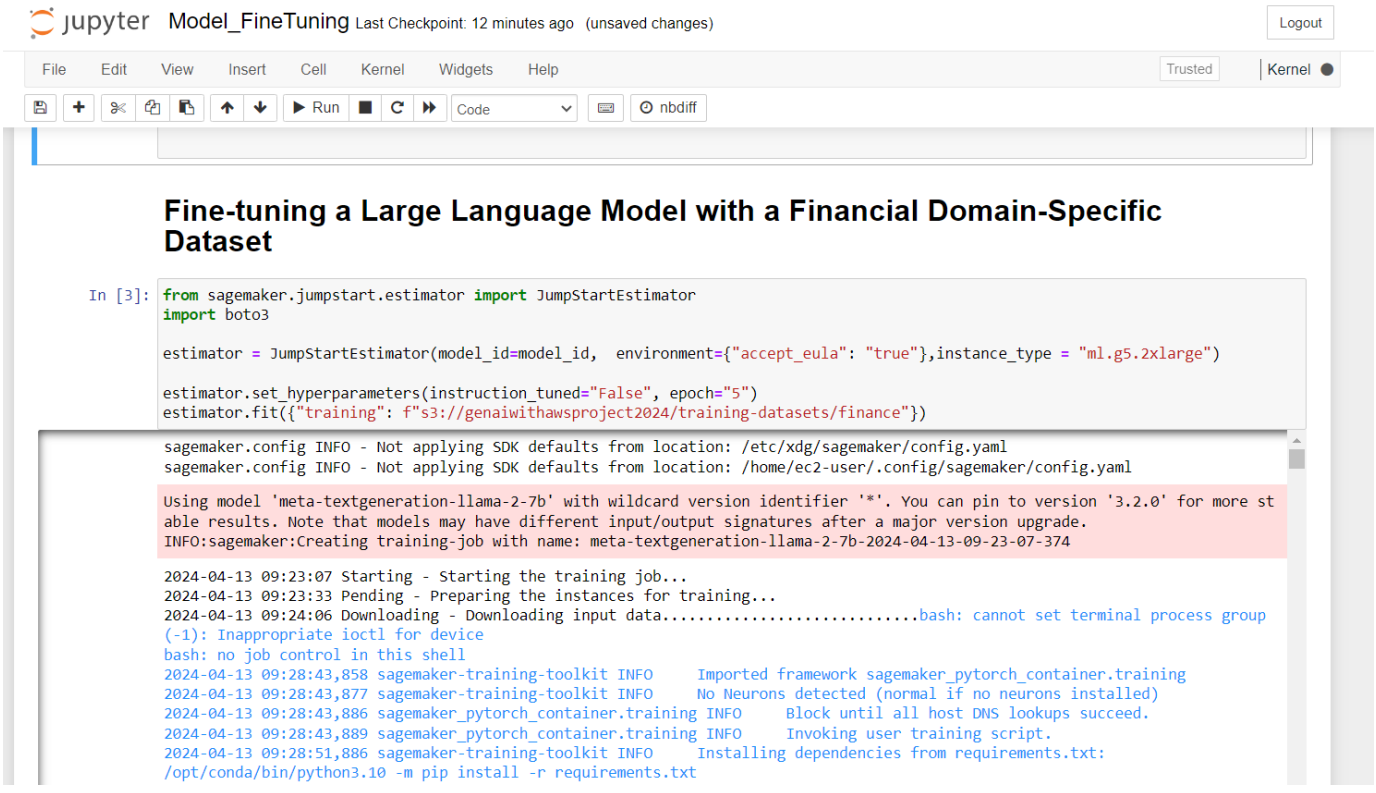

Training Completed

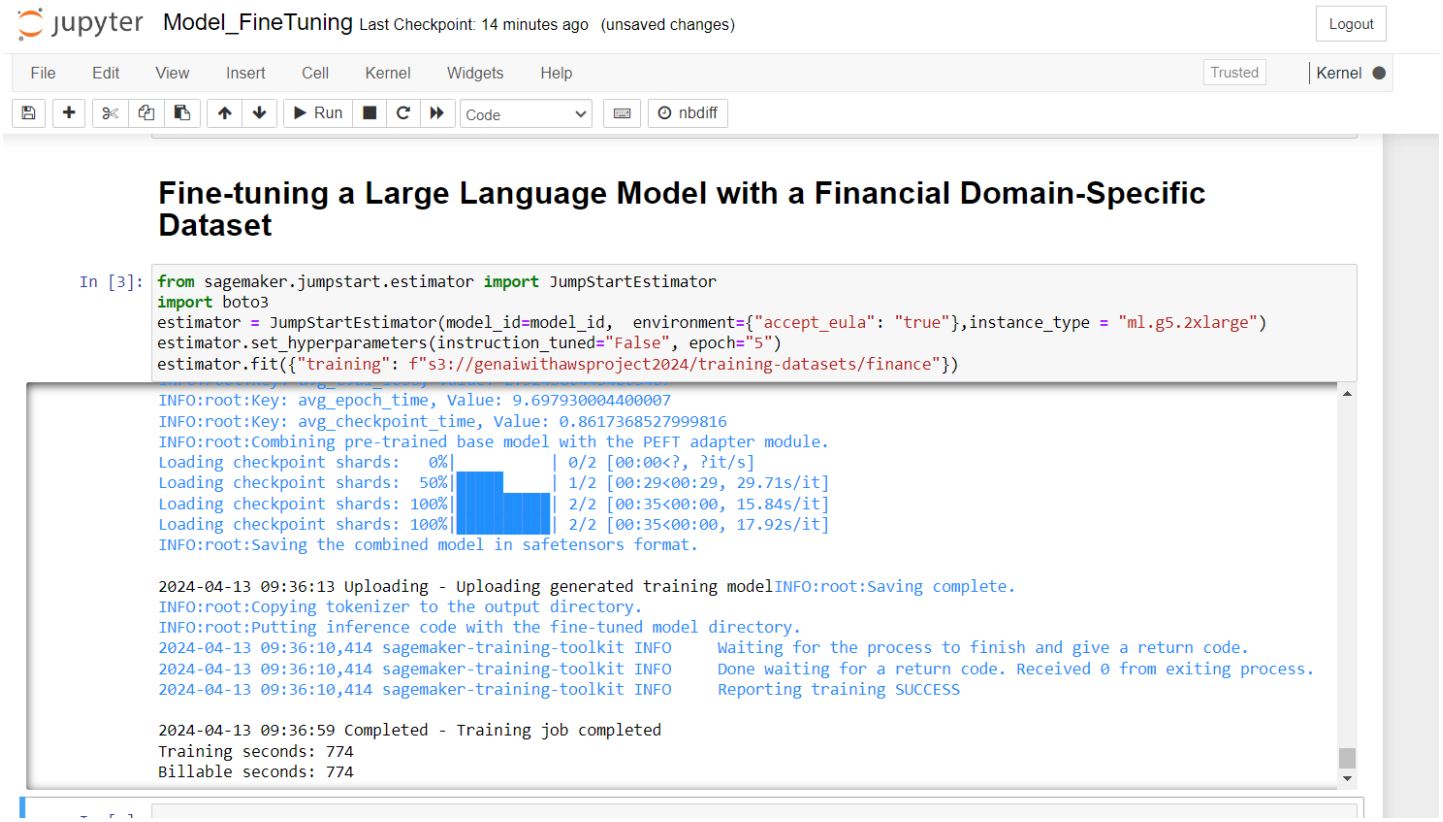

Step 5: Deploying the Fine-tuned Llama2 Model on AWS Sagemaker

After fine-tuning, the fine-tuned Llama2 model was successfully deployed on SageMaker. The deployment process, including the necessary configurations and settings, is documented in the **model_finetuning.ipynb** notebook cell output. Screenshots in the report showcase the deployed fine-tuned model in the SageMaker environment, providing visual evidence of the deployment steps

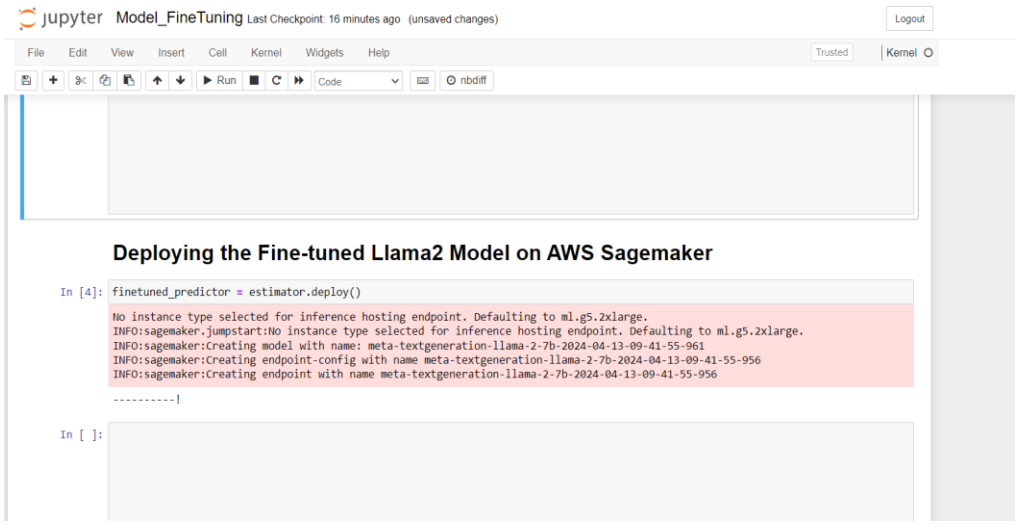

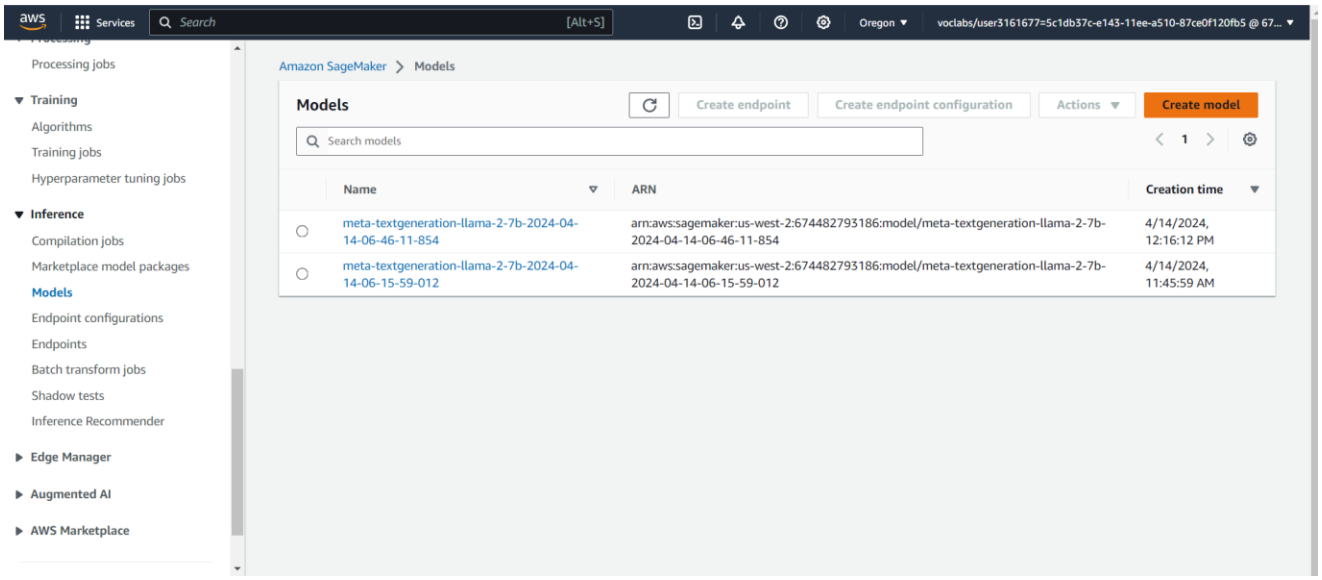

Step 6: Evaluating the Fine-tuned Llama2 Text Generation Large Language Model on Text Generation Tasks and Domain Knowledge

The evaluation of the fine-tuned model focused on assessing its performance in generating domain-specific text. Domain-specific inputs were provided to the fine-tuned model to evaluate its text generation capabilities within the financial domain. The **model_finetuning.ipynb** notebook includes examples of the model's output post-fine-tuning, showcasing improvements or changes in the model's performance. Documentation of these improvements or changes is provided in the Project, highlighting the model's enhanced ability to generate contextually relevant and insightful content within the financial domain.

The Same Inputs are given here:

- "The investment tests performed indicate"
- "the relative volume for the long out of the money options, indicates"
- "The results for the short in the money options"
- "The results are encouraging for aggressive investors"

Examples of the model's output post-fine-tuning:

jupyter Model_FineTuning Last Checkpoint: 20 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted Kernel

Run Code nbdiff

Evaluating the Fine-tuned Llama2 Text Generation Large Language Model on Text Generation Tasks and Domain Knowledge

In [8]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response}")
    print("\n==================================\n")

payload = {
    "inputs": "The results for the short in the money options",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = finetuned_predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
The results for the short in the money options
> [{'generated_text': ' are calculated using the option's premium. The premium is the amount the buyer paid for the option.\nTh
e results for the short in the money options are calculated using the option's premium. The premium is the amount the buyer pai
d for the option. The buyer of an option'}]

==================================
```

jupyter Model_FineTuning Last Checkpoint: 21 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted Kernel

Run Code nbdiff

Evaluating the Fine-tuned Llama2 Text Generation Large Language Model on Text Generation Tasks and Domain Knowledge

In [9]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response}")
    print("\n==================================\n")

payload = {
    "inputs": "The results are encouraging for aggressive investors",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = finetuned_predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
The results are encouraging for aggressive investors
> [{'generated_text': ".\nThe first three months of 2019 were not kind to the stock market. The S&P 500 was down 13.5% from its
record highs in early February. It wasn't the first time this year that stocks had a rough start to the year"}]

==================================
```

jupyter Model_FineTuning Last Checkpoint: 21 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted Kernel

Run Code nbdiff

Evaluating the Fine-tuned Llama2 Text Generation Large Language Model on Text Generation Tasks and Domain Knowledge

In [10]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response}")
    print("\n==================================\n")

payload = {
    "inputs": "the relative volume for the long out of the money options, indicates",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = finetuned_predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
the relative volume for the long out of the money options, indicates
> [{'generated_text': ' a higher volume of put buying relative to call buying.\nThe relative volume is a function of the number of options traded.\nThe relative volume is a function of the number of options trading.\nThe relative volume is a function of the number of options traded.\nThe relative volume is a'}]

==================================
```

jupyter Model_FineTuning Last Checkpoint: 22 minutes ago (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted Kernel

Run Code nbdiff

Evaluating the Fine-tuned Llama2 Text Generation Large Language Model on Text Generation Tasks and Domain Knowledge

In [11]:
```
def print_response(payload, response):
    print(payload["inputs"])
    print(f"> {response}")
    print("\n==================================\n")

payload = {
    "inputs": "Domain specific input chosen from above",
    "parameters": {
        "max_new_tokens": 64,
        "top_p": 0.9,
        "temperature": 0.6,
        "return_full_text": False,
    },
}
try:
    response = finetuned_predictor.predict(payload, custom_attributes="accept_eula=true")
    print_response(payload, response)
except Exception as e:
    print(e)
```

```
Domain specific input chosen from above
> [{'generated_text': ' to be used as the input of the neural network.\nThe neural network with a given architecture.\nThe number of neurons in the hidden layer.\nThe number of neurons in the output layer.\nThe number of training samples.\nThe number of epochs to train the neural network.\nThe number'}]

==================================
```

Pre-training Evaluation:

Our initial assessment of the pre-trained Meta Llama 2 7B model revealed its limitations in generating domain-specific and contextually relevant content. The outputs lacked depth and failed to capture the nuances of the financial domain.

Fine-tuning and Post-training Evaluation:

Through meticulous fine-tuning on a domain-specific financial dataset, we witnessed a remarkable transformation in the model's performance. The post-training evaluation showcased its ability to generate insightful, accurate, and contextually relevant text within the financial domain. The model demonstrated a deeper understanding of financial terminology, trends, and analysis, significantly enhancing its utility for real-world applications.

Key Findings:

1. **Domain-Specific Data:** Fine-tuning with a curated domain-specific dataset is crucial for improving model performance and generating high-quality outputs aligned with the target domain.
2. **Hyperparameter Tuning:** Careful selection and tuning of hyperparameters, such as epoch count and learning rates, play a pivotal role in optimizing model performance during training.
3. **Evaluation Metrics:** Utilizing comprehensive evaluation metrics tailored to the domain, such as accuracy, relevance, and coherence, provides valuable insights into the model's capabilities and areas for improvement.
4. **Iterative Approach:** Adopting an iterative approach to training, evaluation, and fine-tuning facilitates continuous improvement and refinement of the model's performance over time.

Conclusion

The comparative evaluation between pre-training and post-training phases underscores the transformative impact of fine-tuning on model performance within the financial domain. By adhering to best practices such as leveraging domain-specific data, optimizing hyperparameters, and employing rigorous evaluation metrics, we've unlocked the full potential of the Meta Llama 2 7B model, paving the way for advanced applications in financial analysis, decision-making, and information generation.